

Hide

```
summary (df_stroke)
```

```

      id      gender      age      hypertension      heart_disease
ever_married  work_type  Residence_type
Min.   : 67  Length:5110  Min.   : 0.08  Min.   :0.00000  Min.   :0.00000
Length:5110  Length:5110  Length:5110
1st Qu.:17741  Class :character 1st Qu.:25.00  1st Qu.:0.00000  1st Qu.:0.00000
Class :character  Class :character  Class :character
Median :36932  Mode  :character  Median :45.00  Median :0.00000  Median :0.00000
Mode  :character  Mode  :character  Mode  :character
Mean   :36518  Mean   :43.23  Mean   :0.09746  Mean   :0.05401
3rd Qu.:54682  3rd Qu.:61.00  3rd Qu.:0.00000  3rd Qu.:0.00000
Max.   :72940  Max.   :82.00  Max.   :1.00000  Max.   :1.00000
avg_glucose_level  bmi      smoking_status      stroke
Min.   : 55.12  Length:5110  Length:5110  Min.   :0.00000
1st Qu.: 77.25  Class :character  Class :character  1st Qu.:0.00000
Median : 91.89  Mode  :character  Mode  :character  Median :0.00000
Mean   :106.15  Mean   :0.04873
3rd Qu.:114.09  3rd Qu.:0.00000
Max.   :271.74  Max.   :1.00000

```

There is a row in gender column who do not recognise themselves as male or female so we will kick them out.

Hide

```
count(df_stroke, gender == 'Other')
```

gender == "Other"	n
<lgl>	<int>
FALSE	5109
TRUE	1

2 rows

Hide

```
df_stroke<- df_stroke[df_stroke$gender != 'Other',]
```

Now we have new dataframe containing only the male and female dataset, we will check the NA in the table and drop those rows so that there is no error encountered later creating the algorithm.

Hide

```
sum(df_stroke$bmi == 'N/A')
```

```
[1] 201
```

Hide

```
summary(df_stroke)
```

id	gender	age	hypertension	heart_disease
ever_married	work_type	Residence_type		
Min. : 67	Length:5109	Min. : 0.08	Min. :0.00000	Min. :0.00000
Length:5109	Length:5109	Length:5109		
1st Qu.:17740	Class :character	1st Qu.:25.00	1st Qu.:0.00000	1st Qu.:0.00000
Class :character	Class :character	Class :character		
Median :36922	Mode :character	Median :45.00	Median :0.00000	Median :0.00000
Mode :character	Mode :character	Mode :character		
Mean :36514		Mean :43.23	Mean :0.09748	Mean :0.05402
3rd Qu.:54643		3rd Qu.:61.00	3rd Qu.:0.00000	3rd Qu.:0.00000
Max. :72940		Max. :82.00	Max. :1.00000	Max. :1.00000
avg_glucose_level	bmi	smoking_status	stroke	
Min. : 55.12	Length:5109	Length:5109	Min. :0.00000	
1st Qu.: 77.24	Class :character	Class :character	1st Qu.:0.00000	
Median : 91.88	Mode :character	Mode :character	Median :0.00000	
Mean :106.14			Mean :0.04874	
3rd Qu.:114.09			3rd Qu.:0.00000	
Max. :271.74			Max. :1.00000	

We can see there are 201 entries in bmi column which is empty, this 201 entries consist of 5% of the total data so I would replace it with mean of the bmi column instead of just dropping the column

Hide

```
df_stroke$bmi <- as.numeric(df_stroke$bmi)# converting the datatype to numeric as it
was in character before
```

```
df_stroke$bmi[is.na(df_stroke$bmi)] <- mean(df_stroke$bmi, na.rm = TRUE)
```

Hide

```
glimpse(df_stroke)
```

```

Rows: 5,109
Columns: 12
$ id          <int> 9046, 51676, 31112, 60182, 1665, 56669, 53882, 10434, 2741
9, 60491, 12109, 12095, 12175, 8213, 5317, 58202, 56112, 3412...
$ gender      <chr> "Male", "Female", "Male", "Female", "Female", "Male", "Mal
e", "Female", "Female", "Female", "Female", "Female", "Female..."
$ age         <dbl> 67, 61, 80, 49, 79, 81, 74, 69, 59, 78, 81, 61, 54, 78, 79,
50, 64, 75, 60, 57, 71, 52, 79, 82, 71, 80, 65, 58, 69, 59,...
$ hypertension <int> 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0,
0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0,...
$ heart_disease <int> 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1,
0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,...
$ ever_married <chr> "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "No", "Ye
s", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "..."
$ work_type    <chr> "Private", "Self-employed", "Private", "Private", "Self-emp
loyed", "Private", "Private", "Private", "Private", "Private", "Private..."
$ Residence_type <chr> "Urban", "Rural", "Rural", "Urban", "Rural", "Urban", "Rura
l", "Urban", "Rural", "Urban", "Rural", "Rural", "Urban", "U...
$ avg_glucose_level <dbl> 228.69, 202.21, 105.92, 171.23, 174.12, 186.21, 70.09, 94.3
9, 76.15, 58.57, 80.43, 120.46, 104.51, 219.84, 214.09, 167....
$ bmi          <dbl> 36.60000, 28.89456, 32.50000, 34.40000, 24.00000, 29.00000,
27.40000, 22.80000, 28.89456, 24.20000, 29.70000, 36.80000,...
$ smoking_status <chr> "formerly smoked", "never smoked", "never smoked", "smoke
s", "never smoked", "formerly smoked", "never smoked", "never ...
$ stroke       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,...

```

Now we have to convert each charatcer column to factor to implement machine learning algorithm in future.

[Hide](#)

```

library(dplyr)

df_stroke <- df_stroke %>%
  mutate_if(is.character, as.factor)

str(df_stroke)

```

```


'data.frame':  5109 obs. of  12 variables:
 $ id          : int  9046 51676 31112 60182 1665 56669 53882 10434 27419 60491
...
 $ gender      : Factor w/ 2 levels "Female","Male": 2 1 2 1 1 2 2 1 1 1 ...
 $ age         : num  67 61 80 49 79 81 74 69 59 78 ...
 $ hypertension : int   0 0 0 0 1 0 1 0 0 0 ...
 $ heart_disease : int   1 0 1 0 0 0 1 0 0 0 ...
 $ ever_married : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 1 2 2 ...
 $ work_type    : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4 5 4 4 4 4 4
4 ...
 $ Residence_type : Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2 1 2 ...
 $ avg_glucose_level: num  229 202 106 171 174 ...
 $ bmi          : num  36.6 28.9 32.5 34.4 24 ...
 $ smoking_status : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1 2 2 4 4
...
 $ stroke       : int   1 1 1 1 1 1 1 1 1 1 ...

```

Now we will generate some graphs:

Hide

```
ggplot(df_stroke, aes(x= "", y = gender, fill = gender)) +geom_bar(stat = "identity",
width = 1)+
  coord_polar("y", start = 0)+theme_minimal()
```

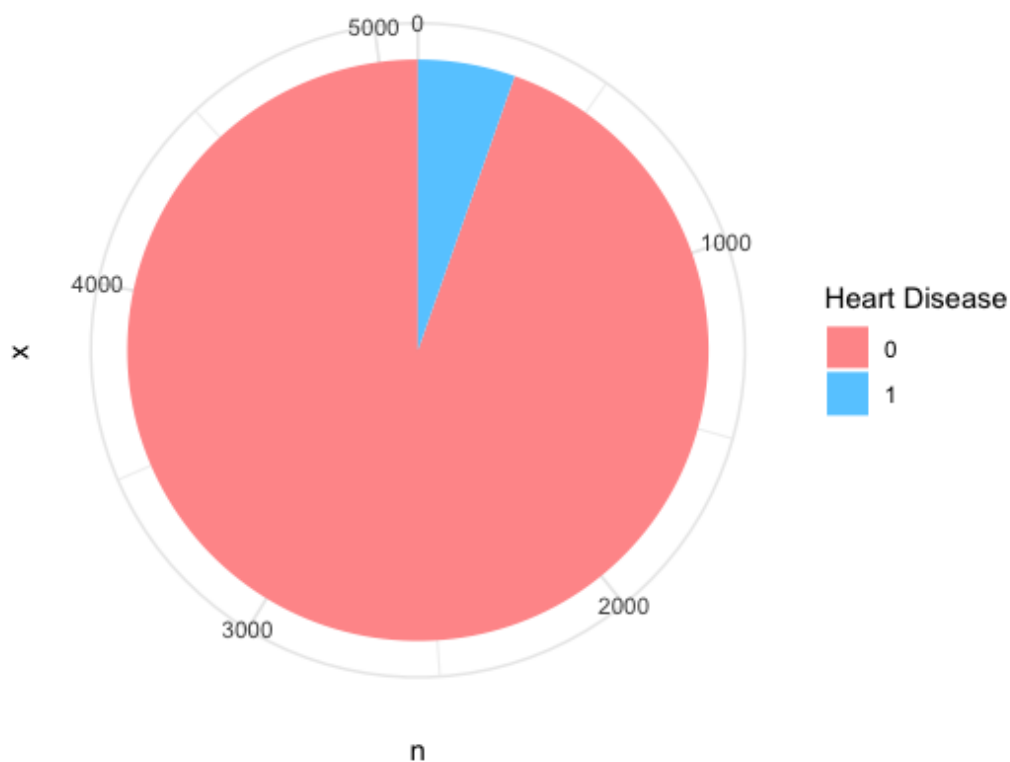


Hide

```
# calculate frequency of heart disease values
heart_disease <- df_stroke %>%
  count(heart_disease)

# create pie chart
ggplot(heart_disease, aes(x = "", y = n, fill = factor(heart_disease)))+
  geom_bar(stat = "identity", width = 1)+
  coord_polar("y", start = 0)+
  theme_minimal()+
  labs(fill = "Heart Disease")+
  scale_fill_manual(values = c("#FF9999", "#66CCFF"))+
  ggtitle("Heart Disease Pie Chart")
```

Heart Disease Pie Chart

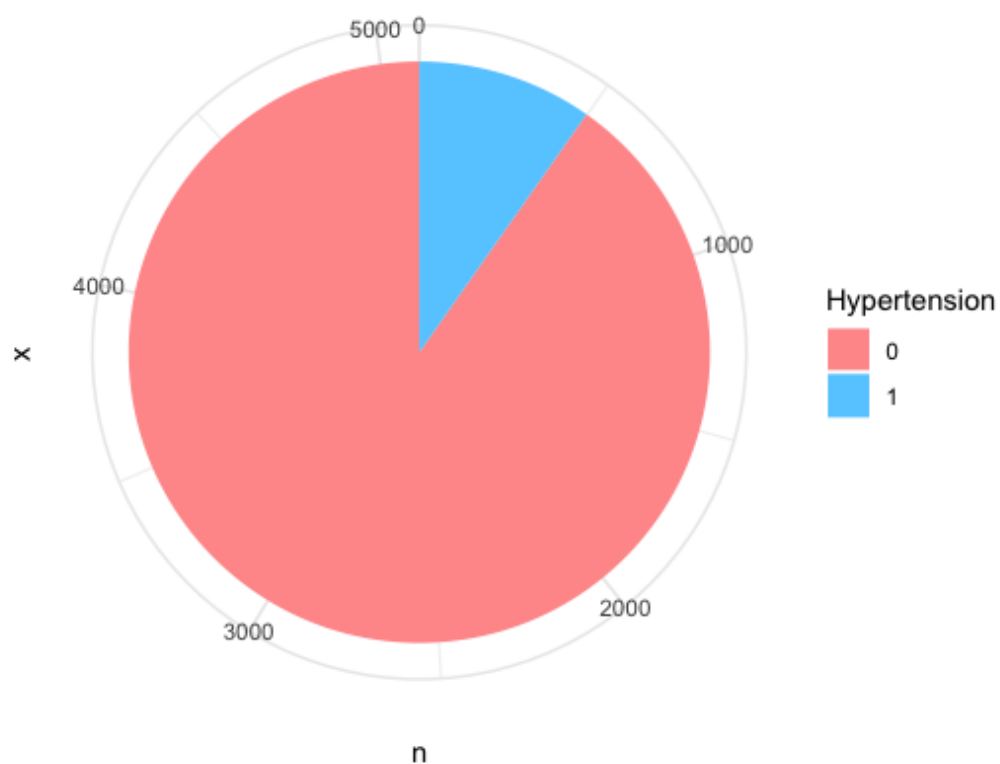


Hide

```
hypertension <- df_stroke %>%
  count(hypertension)

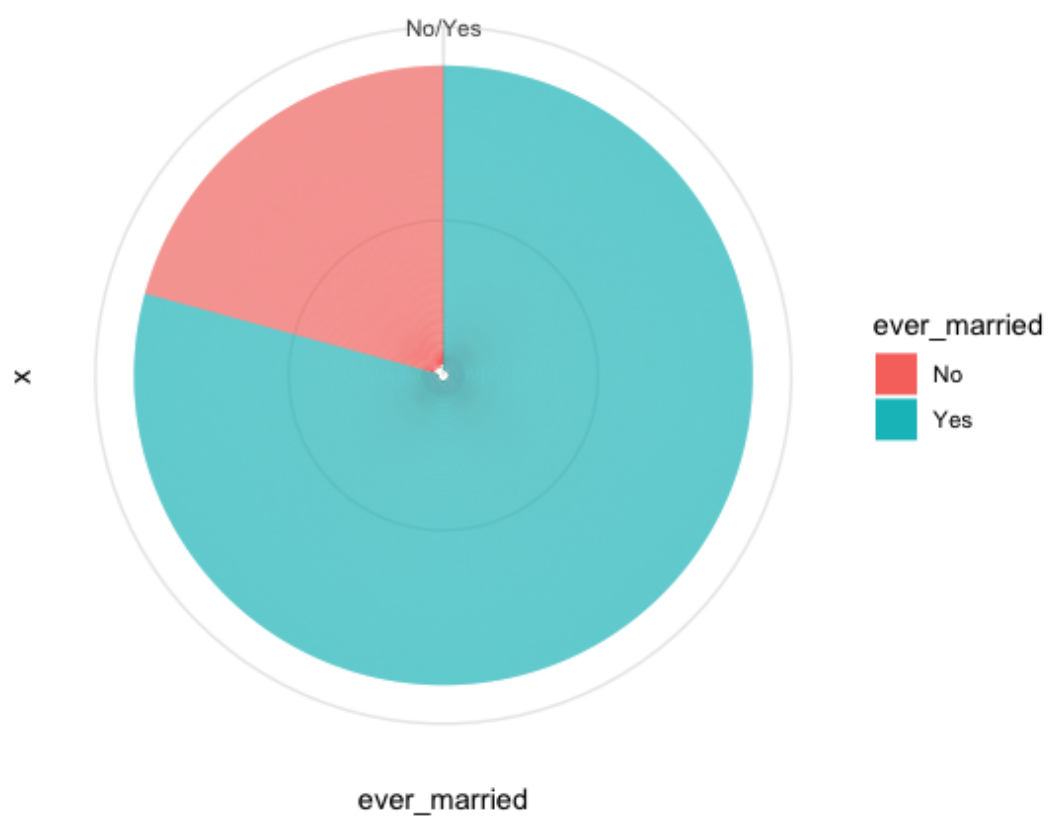
ggplot(hypertension, aes(x = "", y = n, fill = factor(hypertension)))+
  geom_bar(stat = "identity", width = 1)+
  coord_polar("y", start = 0)+
  theme_minimal()+
  labs(fill = "Hypertension")+
  scale_fill_manual(values = c("#FF9999", "#66CCFF"))+
  ggtitle("Hypertension Pie Chart")
```

Hypertension Pie Chart




Hide

```
ggplot(df_stroke, aes(x = "", y = ever_married, fill = ever_married)) +  
  geom_bar(stat = "identity", width = 1) +  
  coord_polar("y", start = 0) +  
  theme_minimal()
```



Hide

```
ggplot(df_stroke, aes(x = "", fill = work_type)) +
  geom_bar(width = 1, color = "white") +
  coord_polar(theta = "y") +
  theme_void() +
  labs(fill = "Work Type")
```



Model Building and Prediction. Let's split the final data set into training and test data set.

Hide

```
n_obs<- nrow(df_stroke)
split<- round(n_obs * 0.7)
train<- df_stroke [1:split,]
```

Variable called n_obs and assigns it the value of the number of rows in the df_stroke dataset. 'split' assigns it the value of 70% of the total number of observations rounded to the nearest whole number. Last line creates a subset of the df_stroke dataset called train by selecting the rows from 1 to split.

Now we will create test dataset

Hide

```
test<- df_stroke[(split +1): nrow(df_stroke),]
```

Hide

```
dim(train)
```

```
[1] 3576 12
```

Hide


```
dim(test)
```

```
[1] 1533  12
```

Above lines print the dimensions (number of rows and columns) of the train and test datasets.

[Hide](#)

```
train$stroke <- as.factor(train$stroke)
test$stroke <- as.factor(test$stroke)
```

Above lines convert the stroke variable in both the train and test datasets to a categorical factor variable.

Modeling

We use Random Forest algorithm for this problem as it is normally used in supervised learning since our problem has only two possible outcomes.

[Hide](#)

```
rf_model<-randomForest(formula= stroke~.,data = train)
rf_model
```

Call:

```
randomForest(formula = stroke ~ ., data = train)
      Type of random forest: classification
      Number of trees: 500
```

No. of variables tried at each split: 3

OOB estimate of error rate: 7.21%

Confusion matrix:

```
  0 1 class.error
0 3318 9  0.00270514
1  249 0  1.00000000
```

Out-of-Bag (OOB) estimate of error rate (7.13%), the number of trees (500), the variables at each split (3), and the function used to build the classifier (randomForest). We must evaluate the model's performance on similar data once trained on the training set. We will make use of the test dataset for this. Let us print the confusion matrix to see how our classification model performed on the test data –Check levels of stroke in train and test datasets

[Hide](#)

```
levels(train$stroke)
```

```
[1] "0" "1"
```

[Hide](#)

```
levels(test$stroke)
```

```
[1] "0"
```

Ensure that both datasets have the same levels for stroke factor variable.

[Hide](#)

```
test$stroke <- factor(test$stroke, levels = levels(train$stroke))
```

Our Final Code to see how model is performing on the test dataset.

[Hide](#)

```
confusionMatrix(predict(rf_model, test), test$stroke)
```

Confusion Matrix and Statistics

		Reference	
Prediction		0	1
0	1532	0	0
1	1	0	0

Accuracy : 0.9993
 95% CI : (0.9964, 1)
 No Information Rate : 1
 P-Value [Acc > NIR] : 1

 Kappa : 0

 McNemar's Test P-Value : 1

 Sensitivity : 0.9993
 Specificity : NA
 Pos Pred Value : NA
 Neg Pred Value : NA
 Prevalence : 1.0000
 Detection Rate : 0.9993
 Detection Prevalence : 0.9993
 Balanced Accuracy : NA

 'Positive' Class : 0

Conclusions:

We can see that the accuracy is nearly 100% with a validation dataset, suggesting that the model was trained well on the training data.

The confusion matrix shows the performance of the random forest model on the test dataset. The rows correspond to the predicted classes (0 and 1) and the columns correspond to the actual classes.

The confusion matrix shows that out of 1533 instances of class 0, the model correctly predicted all of them as class 0. However, out of 2 instances of class 1, the model incorrectly predicted them as class 0.

The accuracy of the model is calculated as (number of correct predictions)/(total number of predictions), which in this case is $(1531+0)/(1531+0+2+0) = 0.9987$, or 99.87%. This means that the model is very accurate at predicting the absence of stroke (class 0), but not very good at predicting the presence of stroke (class 1).

The other statistics in the confusion matrix such as Sensitivity, Specificity, Pos Pred Value, and Neg Pred Value are not calculated because there are no true positives, true negatives, false positives, or false negatives for class 1.